

CONCEPTUAL PAPER / PREPRINT DRAFT

Nash's Cage

Robust Viability-Constrained Payoff-Inversion Governance
under Unknown Irreversible Boundaries

SPIRALREALITY

Conceptual architecture with Ryō ∴ SpiralArchitect

Civilization is optimizing inside an unknown viability boundary.

Regenerated revision v0.2

7 August 2026

This manuscript is a formal research proposal, not a claim of empirical validation or a finished institutional design.

Abstract

Climate governance is commonly framed as a cooperation problem in which individually rational actors underprovide a global public good. That framing is useful but incomplete. It does not, by itself, represent endogenous institutional delay, strategic metric manipulation, symbolic compliance, unequal externalization capacity, uncertain recovery boundaries, or the possibility that governance capacity deteriorates while physical risk accumulates. This paper introduces *Nash's Cage*: a partially observable, capture-prone stochastic dynamic game in which self-interested agents optimize short-horizon payoffs while externalizing physical and social constraints, thereby shrinking a robust viability region faster than institutions can observe, coordinate, and respond.

We formalize a companion architecture, the *Robust Viability-Constrained Cage-Inversion Model* (RVCIM). The model combines heterogeneous agent utilities, a manipulable KPI observation layer, endogenous institutional actuators, narrative dynamics, adversarial capture, model ambiguity, a robust viability kernel, and a braking-time measure termed *controllability reserve*. Four conditional propositions isolate the architecture's central claims: monitoring without institutional coupling is performative; capture can dominate nominal triggers; non-positive braking reserve removes a robust certificate of timely response and, under a separately declared precautionary rule, triggers escalation; and cooperation becomes locally rational only when the payoff differential is inverted for a decisive coalition. Viability and justice are treated as hard constraints rather than fungible terms in a welfare sum.

The paper then specifies accountable automatic triggers, anti-capture separation of powers, multimodal ecological sensing, and dynamic policy pathways. Soundscape metrics are included only as one candidate sensor within a multimodal ensemble requiring site-specific validation. A red-team catalogue describes realistic adaptations by regulated actors and corresponding defenses. Finally, we state falsifiable hypotheses, a simulation program, and an executable F0 reference contract. The contribution is not a proof that a particular climate treaty must fail, nor a claim that one universal controller can govern Earth. It is a formal specification for asking whether a governance system preserves the capacity to brake before unknown irreversible boundaries become operationally unreachable.

Keywords: dynamic games; climate governance; viability theory; robust control; institutional capture; tipping points; deep uncertainty; control barrier functions; soundscape ecology; just transition.

Claim boundary

The term *Nash's Cage* does not imply that every Nash equilibrium is socially harmful, that climate politics reduces to a two-player Prisoner's Dilemma, or that automation can replace democratic judgment. It names a narrower model class: stable strategic configurations in which local best responses reproduce collective physical degradation and degrade the future capacity to change course.

Revision provenance

The repository preserves the exact uploaded v0.1 source. This v0.2 is a regenerated revision dated 7 August 2026, based on that source and extended with the executable-companion contract described below. It is not represented as byte-identical to an earlier, unavailable v0.2 build. File identities and the regeneration boundary are recorded in the repository's release manifest.

Contents

1	Introduction	3
2	Foundations and Scope	3
2.1	From bad equilibrium to a dynamic cage	4
2.2	Partial observability and endogenous institutions	4
2.3	Viability, safety, and deep uncertainty	4

2.4	Earth-system boundaries and irreversible risk	4
3	Model Primitives	5
3.1	State, agents, and action space	5
3.2	Utility, horizons, and externalization	6
3.3	Observation, KPI construction, and manipulation	6
3.4	Institutional and narrative dynamics	6
4	Nash's Cage: Definition and Local Results	7
4.1	Cooperation, defection, and decisive actors	7
4.2	Conditional propositions	8
5	Unknown Irreversible Boundaries	8
5.1	Safety is not viability	8
5.2	Functional irreversibility	9
5.3	Controllability reserve	9
5.4	Endogenous optimism and doom	10
6	The RVCIM Governance Architecture	10
6.1	Viability and justice as hard constraints	10
6.2	Robust barrier conditions	10
6.3	Payoff inversion controller	11
6.4	Trigger modes	11
6.5	One-way doors and burden of proof	11
6.6	Accountable automation and institutional separation	12
7	Observation Architecture and Ecological Sensing	12
7.1	A dashboard of non-fungible signals	12
7.2	Soundscape as one ensemble sensor	12
8	Red-Team: How the New OS Will Be Gamed	13
9	Empirical Anchors and Claim Boundaries	15
10	A Falsifiable Research Program	16
10.1	Minimal simulation environment	16
10.2	Hypotheses	16
10.3	Stress tests	17
11	Executable Reference and Feasibility Contract	17
11.1	Feasibility ladder and admissible claims	17
11.2	Paper-to-code correspondence	18
11.3	Reference implementation contract	18
12	Limitations and Open Problems	19
13	Conclusion	20
A	Notation Summary	20
B	Minimal Control Loop	21
C	Cage-Inversion Constitutional Invariants	22
	References	22

1 Introduction

A Nash equilibrium is a profile from which no player benefits by unilateral deviation, given the strategies of the others (Nash, 1950). Stability in this sense is not equivalent to collective optimality. Common-pool resource problems, public-goods games, repeated dilemmas, and climate cooperation studies show that individually defensible choices can maintain collectively inferior outcomes (Hardin, 1968; Axelrod, 1984; Ostrom, 1990; Milinski et al., 2008; Nordhaus, 2015). Yet the climate problem is more difficult than the canonical matrix game: actors are heterogeneous, actions are continuous and path dependent, observations are delayed, institutional rules change endogenously, and some actors can spend resources to reshape the rules by which they are evaluated.

The empirical context motivates, but does not prove, the model. The Paris Agreement's Article 15 mechanism is expressly expert-based, facilitative, non-adversarial, and non-punitive (United Nations Framework Convention on Climate Change, 2015; United Nations Framework Convention on Climate Change, 2026). UNEP's 2025 assessment estimated end-of-century warming of roughly 2.3–2.5 °C under full implementation of current nationally determined contributions and 2.8 °C under current policies (United Nations Environment Programme, 2025). The 2025 Production Gap Report found that planned 2030 fossil-fuel production remained more than twice the level associated with a 1.5 °C pathway (Stockholm Environment Institute et al., 2025). The WMO reported a global mean atmospheric carbon dioxide concentration of 423.9 ppm in 2024 (World Meteorological Organization, 2025). These observations are consistent with a gap between symbolic commitment and physical trajectory, but they do not by themselves identify a single causal mechanism. RVCIM is proposed as a testable mechanism class.

Core claim

Diagnostic thesis. The central governance failure is not merely insufficient observation. It is weak, delayed, and manipulable coupling from physical observation to institutional response, from institutional response to agent payoffs, and from payoffs to actions that preserve future controllability.

The paper makes six contributions.

1. It defines Nash's Cage as a partially observable, institutionally endogenous dynamic game rather than a static metaphor.
2. It distinguishes physical safety, robust viability, and operational braking reserve, allowing a state to appear safe while already lacking a certified path to stop.
3. It defines *payoff coupling strength*, linking KPI changes to the utility difference between cooperation and defection.
4. It models regulatory capture as an adversarial action vector that attacks sensors, models, actuators, delay, narratives, finance, and justice.
5. It specifies RVCIM, in which viability and justice are hard constraints, automatic triggers are accountable rather than discretionary, and policy pathways remain adaptive under deep uncertainty.
6. It converts the architecture into a red-team and empirical research program with measurable failure criteria.

2 Foundations and Scope

2.1 From bad equilibrium to a dynamic cage

The Prisoner's Dilemma is a useful local motif: if temptation exceeds mutual cooperation, mutual defection, and unilateral cooperation in the familiar ordering, defection is individually dominant while mutual cooperation is Pareto superior. Climate interactions, however, can also resemble assurance games, chicken, repeated public-goods games, bargaining games, common-pool resource games, and coalition games. Accordingly, Nash's Cage does not assert one universal payoff matrix. It asserts a dynamic property: for a decisive set of actors, the locally preferred actions reproduce degradation and resist unilateral exit.

Repeated interaction can sustain cooperation when future consequences, reputation, reciprocity, and credible punishment matter (Axelrod, 1984; Ostrom, 1990). Climate governance often weakens this "shadow of the future" through election cycles, quarterly reporting, capital turnover, jurisdictional leakage, and non-binding commitments. Climate clubs and coordinated trade measures have therefore been proposed as methods for changing free-riding incentives (Nordhaus, 2015). RVCIM generalizes the same intuition: the object of design is not moral exhortation but the intertemporal payoff field.

2.2 Partial observability and endogenous institutions

Dynamic noncooperative games explicitly represent evolving states and strategic feedback (Başar and Olsder, 1999). Partially observable decision processes distinguish the true state from the information available to the decision-maker (Kaelbling et al., 1998). In planetary governance, this distinction is essential: no actor observes the complete Earth-system state, the full distribution of future thresholds, or the true degree of institutional fragility. Moreover, the observation process is not neutral. Reporting boundaries, measurement choices, lobbying, accounting rules, and narrative framing can be strategic actions.

Institutions are therefore modeled as endogenous state variables and actuators, not fixed constraints. Regulatory capture research emphasizes that regulated interests can influence agencies, information, standards, and enforcement (Carpenter and Moss, 2014). Greenwashing research similarly treats symbolic environmental communication as an organizational response shaped by incentives and monitoring conditions (Delmas and Burbano, 2011). In the present model, symbolic action and capture are explicit elements of the action space.

2.3 Viability, safety, and deep uncertainty

Viability theory studies whether a controlled dynamical system can remain within a constraint set over time (Aubin et al., 2011). Control barrier functions encode safety conditions that can be enforced while pursuing performance objectives (Ames et al., 2017; Ames et al., 2019). Robust and adaptive decision frameworks address uncertainty in models, distributions, and future conditions; dynamic adaptive policy pathways emphasize sequencing, signposts, and switching rules rather than commitment to a single forecast (Haasnoot et al., 2013; Kwakkel et al., 2016).

The analogy to social governance must be used carefully. Institutions are not deterministic motors, citizens are not control inputs, and political legitimacy cannot be reduced to a quadratic program. The transfer is structural: safety should constrain optimization; uncertain boundaries require robustness; delay belongs inside the model; and a system should preserve the option to change course.

2.4 Earth-system boundaries and irreversible risk

Earth-system and planetary-boundary research identifies interacting biophysical processes that delimit a safe operating space (Rockström et al., 2023; Richardson et al., 2023). Climate tipping elements may exhibit threshold behavior, self-amplifying change, and hysteresis (Scheffer et al., 2009; Armstrong McKay et al., 2022). The IPCC assesses increasingly severe, interconnected, and in some cases irreversible risks with additional warming (Intergovernmental Panel on Climate Change, 2022; Intergovernmental Panel

on Climate Change, 2023). These literatures do not provide a single exact global cliff. They instead motivate the central uncertainty of this paper: the shape and distance of the true recovery boundary are only partially known.

3 Model Primitives

3.1 State, agents, and action space

Let time be discrete, $t = 0, 1, \dots$, and let $\mathcal{N} = \{1, \dots, N\}$ denote heterogeneous actors. An actor may represent a state, firm, financial institution, regulator, local government, sectoral coalition, or organized public. The full state is

$$\mathbf{x}_t = (s_t, \mathbf{z}_t, \mathbf{u}_t, \mathbf{n}_t), \quad (3.1)$$

where s_t is the physical state, $\mathbf{z}_t$ contains actor-specific stocks and dependencies, $\mathbf{u}_t$ is the institutional state, and $\mathbf{n}_t$ is the cognitive and legitimacy state.

Table 1: Layered state representation. The variables are illustrative rather than exhaustive.

Layer	Representative variables	Interpretation
Physical	C_t, B_t, X_t, L_t, W_t	Atmospheric forcing; biosphere integrity; tail-event intensity; infrastructure stress; water, food, and coupled life-support states.
Agent	$E_{i,t}, F_{i,t}, D_{i,t}, G_{i,t}$	Emissions; fossil dependence or production; direct decarbonization; symbolic compliance.
Institution	$P_t, Q_t, M_t, \tau_t, J_t$	Penalty strength; transition support; monitoring fidelity; policy delay; justice and compensation capacity.
Cognition	$\Phi_t, A_t^{\text{ware}}, T_t^{\text{rust}}$	Dominant narrative frame; public awareness; trust and legitimacy.
Irreversibility	$\mathcal{V}_t, CR_t, \rho_t$	Viability region; controllability reserve; robust recovery probability.

Actor i chooses

$$\mathbf{a}_{i,t} = (e_{i,t}, f_{i,t}, d_{i,t}, g_{i,t}, \ell_{i,t}, r_{i,t}, h_{i,t}), \quad (3.2)$$

where the components denote emissions, fossil production or dependence, direct decarbonization, symbolic action, capture or delay effort, ecological restoration, and adaptation. Including $g_{i,t}$ and $\ell_{i,t}$ is essential: a model that treats performative compliance and rule-shaping as exogenous noise omits strategies that may be cheaper than physical transformation.

The physical dynamics are

$$s_{t+1} = F_\theta(s_t, \mathbf{a}_t, \xi_t), \quad \theta \in \Theta_t, \quad (3.3)$$

with aggregate action $\mathbf{a}_t = (\mathbf{a}_{1,t}, \dots, \mathbf{a}_{N,t})$, disturbance ξ_t , and model index θ . Illustrative components are

$$C_{t+1} = C_t + \sum_i e_{i,t} - \text{Sink}_\theta(B_t, C_t, \text{heat, drought, fire}), \quad (3.4)$$

$$B_{t+1} = B_t - \alpha C_t - \beta U_t^{\text{land}} - \gamma U_t^{\text{pollution}} - \delta U_t^{\text{noise/light}} + \sum_i r_{i,t}. \quad (3.5)$$

The sink is endogenous: degradation can reduce future absorptive capacity. Tail risk is represented as

$$q_{j,t}^{\text{tail}} = \mathbb{P}_\theta \left(X_{j,t} > \Theta_j^{\text{design}} \mid s_t \right), \quad (3.6)$$

not merely by a change in the mean. Cascades may be modeled through an interdependency graph among power, water, food, health, logistics, communications, finance, and governance.

3.2 Utility, horizons, and externalization

Let Π_i^{feas} denote actor i 's feasible policies. A generic discounted objective is

$$\begin{aligned} U_i(\pi_i, \pi_{-i}; \theta) = \mathbb{E}_\theta \sum_{k=0}^{H_i} \beta_i^k & \left[G_i^{\text{econ}}(e_i, f_i) + B_i^{\text{free}}(\mathbf{a}_{-i}) + B_i^{\text{sym}}(g_i, \mathbf{n}) + Q_i(d_i, r_i; \boldsymbol{\iota}, \mathbf{k}) \right. \\ & - C_i^{\text{trans}}(d_i, r_i, h_i) - P_i(e_i, f_i; \boldsymbol{\iota}, \mathbf{k}) - D_i(s) \\ & \left. - C_i^{\text{audit}}(g_i, \ell_i; M) - C_i^{\text{legal}}(\mathbf{a}_i; \boldsymbol{\iota}) - C_i^{\text{cap}}(\ell_i) \right]. \end{aligned} \quad (3.7)$$

Short political and financial horizons are represented by small H_i or low β_i . Climate damage borne by the actor is

$$D_i(s_t) = E_i^{\text{exposure}} V_i^{\text{vulnerability}} \eta_i^{\text{internalization}} D(s_t), \quad \chi_i^{\text{externalization}} = 1 - \eta_i^{\text{internalization}}. \quad (3.8)$$

Powerful actors may lower $\eta_i^{\text{internalization}}$ through insurance, relocation, price pass-through, legal privilege, supply-chain displacement, or temporal deferral. Benefits are internalized; losses are displaced across jurisdiction, class, species, or generation.

3.3 Observation, KPI construction, and manipulation

The true state is not directly observed. Raw observations satisfy

$$\mathbf{y}_t = O_\theta(\mathbf{x}_t, \mathbf{a}_t) + \boldsymbol{\varepsilon}_t + \boldsymbol{\mu}_t, \quad (3.9)$$

where $\boldsymbol{\varepsilon}_t$ is measurement and model error and $\boldsymbol{\mu}_t$ is strategic manipulation. The institutional KPI vector is produced by a fusion operator

$$\mathbf{k}_t = \mathcal{H}(\mathbf{y}_{0:t}; \mathcal{M}_t), \quad (3.10)$$

where $\mathcal{M}_t$ contains models, baselines, accounting rules, and metadata. A stylized integrity score is

$$\text{DI}_t = \frac{M_t^{\text{fidelity}} M_t^{\text{independence}} M_t^{\text{comparability}}}{1 + O_t^{\text{opacity}} + U_t^{\text{manipulability}} + F_t^{\text{political}}}. \quad (3.11)$$

More data do not necessarily imply more truth. A sufficiently opaque or incomparable reporting regime can increase audit cost while preserving strategic ambiguity.

3.4 Institutional and narrative dynamics

Let Γ map KPI deviations into pre-legislated adjustments to taxes, subsidies, supply caps, permits, trade rules, public finance, restoration, and compensation. Institutional dynamics are

$$\boldsymbol{\iota}_{t+1} = \text{Proj}_{\mathfrak{Z}} \left[\boldsymbol{\iota}_t + \Gamma(\mathbf{k}_t - \bar{\mathbf{k}}_t) - \mathbf{c}_t - \boldsymbol{\varphi}_t \right], \quad (3.12)$$

where $\mathbf{c}_t$ is strategic capture and $\boldsymbol{\varphi}_t$ is ordinary friction or capacity loss. Capture is decomposed as

$$\mathbf{c}_t = (c_t^K, c_t^I, c_t^\tau, c_t^N, c_t^J, c_t^F), \quad (3.13)$$

attacking the sensor/KPI layer, institutional actuator, response time, narrative, justice reserve, and financial transmission.

Narratives evolve according to

$$\mathbf{n}_{t+1} = \Psi(\mathbf{n}_t, \mathbf{k}_t, \text{events}_t, \text{media}_t, \text{interests}_t, \text{lived impacts}_t). \quad (3.14)$$

Actors optimize against perceived states

$$\tilde{\mathbf{x}}_{i,t} = \mathcal{P}_i(\mathbf{y}_{0:t}, \mathbf{k}_t, \mathbf{n}_t, \text{interests}_i), \quad \mathbf{a}_{i,t} \in \arg \max_{\mathbf{a}_i} U_i(\mathbf{a}_i \mid \tilde{\mathbf{x}}_{i,t}, \mathbf{t}_t). \quad (3.15)$$

Physical reality is therefore mediated by observation and meaning before it becomes action.

4 Nash's Cage: Definition and Local Results

4.1 Cooperation, defection, and decisive actors

For a specified governance problem, partition feasible policies into cooperative and defective sets, Π_i^C and Π_i^D . The partition is empirical and context-dependent: direct verified emissions reduction may be cooperative, while boundary laundering may be defective even when legally compliant. Define

$$\Delta_i(\mathbf{x}_t) = \sup_{\pi_i \in \Pi_i^C} U_i(\pi_i, \pi_{-i}) - \sup_{\pi_i \in \Pi_i^D} U_i(\pi_i, \pi_{-i}). \quad (4.1)$$

Thus $\Delta_i < 0$ means that defection is preferred. A set $\mathcal{D} \subseteq \mathcal{N}$ is *decisive* if its joint control over emissions, production, finance, technology, enforcement, or demand exceeds a problem-specific critical weight.

Orient KPI components so that positive deviations represent worsening conditions, and let $\mathbf{d}_t \in \mathbb{R}_+^m$ be a normalized policy-relevant deterioration direction. The directional coupling from observation to incentives is

$$\Lambda_i(t; \mathbf{d}_t) = D_{\mathbf{d}_t} \Delta_i = (\nabla_{\mathbf{k}} \Delta_i)^\top \mathbf{d}_t \quad (4.2)$$

$$= (\nabla_{\mathbf{t}} \Delta_i)^\top \frac{\partial \mathbf{t}}{\partial \mathbf{k}} \mathbf{d}_t + (\nabla_{\mathbf{n}} \Delta_i)^\top \frac{\partial \mathbf{n}}{\partial \mathbf{k}} \mathbf{d}_t. \quad (4.3)$$

The first term is formal policy transmission; the second captures reputational and legitimacy pathways. We write Λ_i when the direction is clear. An observation system with high fidelity but $\Lambda_i \approx 0$ can document failure without changing it.

Definition 4.1 (Nash's Cage). Over a window $[t, t + H]$, a system is in a Nash's Cage with respect to a decisive set $\mathcal{D}$ when the following conditions jointly hold to a policy-relevant degree:

- (C1) $\Delta_i < 0$ for all or a controlling share of $i \in \mathcal{D}$;
- (C2) a robust viability margin is non-increasing or deteriorating;
- (C3) payoff coupling Λ_i is weak, delayed, or strategically neutralized;
- (C4) effective capture is at least as strong as the institutional correction it opposes; and
- (C5) symbolic compliance grows faster than verified physical improvement.

A simple performative gap is

$$\text{PG}_t = \|\mathbf{g}_t\| - \kappa \left\| \Delta \mathbf{s}_t^{\text{verified},+} \right\|, \quad (4.4)$$

where $\Delta s^{\text{verified},+}$ denotes verified improvement in physical states and κ maps units. Its purpose is diagnostic, not moral: it asks whether communication and accounting are decoupling from physical response.

4.2 Conditional propositions

Proposition 4.2 (Performative monitoring). *Suppose, locally, that Δ_i depends on the KPI only through institutional and narrative channels, all maps are differentiable, and*

$$\frac{\partial \mathbf{u}}{\partial \mathbf{k}} = 0, \quad \frac{\partial \mathbf{n}}{\partial \mathbf{k}} = 0.$$

Then $\Delta_i(t; \mathbf{d}) = 0$ for every direction $\mathbf{d}$. Increasing observation fidelity alone cannot locally change the actor's cooperation–defection payoff differential.

Proof. Apply the chain rule to Equation (4.3). Both transmission terms vanish. The conclusion is local and conditional; new information may still matter if it changes beliefs, narratives, litigation, or institutional design through a path excluded by the assumptions. $\square$

Proposition 4.3 (Capture dominance). *Let $\mathbf{v}$ be an enforcement direction in institutional state space. If*

$$\langle \mathbf{v}, \Gamma(\mathbf{k}_t - \bar{\mathbf{k}}_t) \rangle \leq \langle \mathbf{v}, \mathbf{c}_t + \boldsymbol{\varphi}_t \rangle, \quad (4.5)$$

then the institutional update in that direction is non-positive before projection. A worsening KPI need not strengthen enforcement.

Proof. Take the inner product of Equation (3.12) before projection with $\mathbf{v}$. The net directional increment is the left-hand trigger minus capture and friction. Under Equation (4.5), it is non-positive. $\square$

Proposition 4.4 (Payoff inversion). *If, for a decisive coalition $\mathcal{D}$, $\Delta_i > 0$ under the feasible institutional state and actors select best responses over $\Pi_i^C \cup \Pi_i^D$, then cooperation is locally preferred by each $i \in \mathcal{D}$. This condition is necessary for the intended cage inversion but is not sufficient for physical viability.*

Proof. The first statement follows from the definition in Equation (4.1). Sufficiency fails because physical lags, coordination failures outside $\mathcal{D}$, model error, and already-negative controllability reserve may prevent recovery even after incentives change. $\square$

5 Unknown Irreversible Boundaries

5.1 Safety is not viability

Let $\mathcal{X}_{\text{safe}} \subseteq \mathcal{X}$ be a constraint set: states that satisfy specified physical, social, and justice conditions at the present time. For a known model θ , the viability kernel is

$$\mathcal{V}_\theta = \{ \mathbf{x}_0 \in \mathcal{X}_{\text{safe}} : \exists \boldsymbol{\pi} \in \Pi^{\text{feas}} \text{ such that } \mathbf{x}_t \in \mathcal{X}_{\text{safe}} \forall t \geq 0 \}. \quad (5.1)$$

A state can lie in $\mathcal{X}_{\text{safe}}$ yet outside $\mathcal{V}_\theta$: current conditions have not violated the constraint, but no feasible future policy can prevent eventual exit. This distinction replaces the image of one visible “point of no return” with a control-dependent boundary.

The true model is unknown. Let Θ_t be an ambiguity set containing plausible physical, technological, economic, and institutional dynamics, plus an explicit structural-error allowance. For a measurable set $\mathcal{A}$, define the first-passage delay from decision time t by

$$\tau_{\mathcal{A}}^{(t)} = \inf \{ h \in \mathbb{N}_0 : \mathbf{x}_{t+h} \in \mathcal{A} \}, \quad \inf \emptyset = +\infty. \quad (5.2)$$

Define the robust viability set over horizon H as

$$\mathcal{V}_t^{\text{rob}} = \left\{ \mathbf{x}_t \in \mathcal{X}_{\text{safe}} : \exists \boldsymbol{\pi} \in \Pi^{\text{feas}} \text{ such that } \sup_{\theta \in \Theta_t} \mathbb{P}_{\theta} \left(\min\{\tau_{\mathcal{X}_{\text{safe}}}^{(t)}, \tau_{\mathcal{X}_{\text{irr}}}^{(t)}\} \leq H \mid \mathbf{x}_t, \boldsymbol{\pi} \right) \leq \varepsilon \right\}. \quad (5.3)$$

The same admissible policy must keep the probability of leaving the safe set or entering a policy-defined functionally irreversible region below the risk tolerance across the ambiguity set.

5.2 Functional irreversibility

Irreversibility is separated into four classes:

1. **Physical:** recovery is impossible on human policy timescales.
2. **Hysteretic:** the return threshold differs from the collapse threshold.
3. **Practical:** recovery is physically possible but exceeds feasible time, energy, finance, or infrastructure.
4. **Institutional:** recovery is physically possible but governance, trust, coordination, or legal capacity is insufficient.

Thus

$$\mathcal{X}_{\text{irr}} = \mathcal{X}_{\text{phys}} \cup \mathcal{X}_{\text{hyst}} \cup \mathcal{X}_{\text{practical}} \cup \mathcal{X}_{\text{institutional}}. \quad (5.4)$$

A robust recovery score is

$$\rho_t(H) = \sup_{\boldsymbol{\pi} \in \Pi^{\text{feas}}} \inf_{\theta \in \Theta_t} \mathbb{P}_{\theta} \left(\mathbf{x}_{t+H} \in \mathcal{X}_{\text{safe}}, \tau_{\mathcal{X}_{\text{irr}}}^{(t)} > H \mid \mathbf{x}_t, \boldsymbol{\pi} \right). \quad (5.5)$$

If $\rho_t(H) < \rho_{\min}$, the system is functionally irreversible relative to the chosen horizon, capabilities, ambiguity set, and safety definition. This is not a metaphysical claim that no future civilization could reverse the state; it is an operational statement about feasible control now.

5.3 Controllability reserve

The critical event may occur before a visible boundary is crossed: the system loses enough time that even a correct response cannot take effect before exit. Let

$$T_{\text{exit}}^L(t) = \inf_{\theta \in \Theta_t} Q_{q_L}^{\theta} \left[\tau_{\mathcal{X}_{\text{irr}}}^{(t)} \mid \mathbf{x}_t \right] \quad (5.6)$$

and

$$T_{\text{response}}^U(t) = \sup_{\theta \in \Theta_t} Q_{q_U}^{\theta} \left[T_{\text{detect}} + T_{\text{verify}} + T_{\text{political}} + T_{\text{legal}} + T_{\text{finance}} + T_{\text{infrastructure}} + T_{\text{biophysical}} \right]. \quad (5.7)$$

The controllability reserve is

$$CR_t = T_{\text{exit}}^L(t) - T_{\text{response}}^U(t). \quad (5.8)$$

The quantiles q_L and q_U encode a chosen risk posture; they are not universal constants.

Proposition 5.1 (Delay-induced loss of certified braking capacity). *If $CR_t \leq 0$ and no admissible faster response lies outside the modeled response set, then the system lacks a robust certificate that intervention can take effect before entry into $\mathcal{X}_{\text{irr}}$. Under a separately declared precautionary decision rule, this condition triggers precautionary or emergency escalation even if the observed state remains inside $\mathcal{X}_{\text{safe}}$.*

Proof. By Equation (5.8), the conservative time to exit is no greater than the conservative response time. Therefore the model cannot certify that the response completes first. The timing result establishes loss of certification; escalation follows from the declared precautionary rule, not from the inequality alone. Neither establishes that irreversible entry will occur in the realized world. $\square$

Reserve is multidimensional. Let normalized components represent physical, observability, actuation, coordination, and justice/legitimacy capacity:

$$CR_t^{\text{sys}} = \min\{\tilde{R}_t^P, \tilde{R}_t^O, \tilde{R}_t^A, \tilde{R}_t^C, \tilde{R}_t^J\}. \quad (5.9)$$

The minimum, rather than an average, reflects a bottleneck claim: excellent sensing does not compensate for absent enforcement; strong enforcement does not compensate for zero legitimacy; legitimacy does not restore an already exhausted physical reserve.

5.4 Endogenous optimism and doom

Civil systems often behave as if $\rho_t > 0$ because finance, investment, insurance, electoral legitimacy, and consumption depend on an operative future. Let an optimism distortion satisfy

$$\hat{\rho}_t = \rho_t + o_t + \epsilon_t^P, \quad (5.10)$$

where o_t may be psychological, institutional, or strategic. Optimism is not the only dangerous distortion. Certainty of doom can shorten horizons, reduce cooperation, and produce disengagement. RVCIM therefore adopts neither “recovery is assured” nor “collapse is certain.” It treats uncertainty about high-consequence irreversibility as a reason to preserve options and braking capacity.

6 The RVCIM Governance Architecture

6.1 Viability and justice as hard constraints

A conventional weighted welfare objective can make an irreplaceable loss appear compensable by sufficiently large economic benefit. RVCIM instead uses constrained or lexicographic optimization:

$$\pi^* = \text{lexmin}_{\pi} (\mathcal{R}_{\text{irr}}(\pi), \mathcal{R}_{\text{justice}}(\pi), -\mathcal{W}_{\text{durable}}(\pi)), \quad (6.1)$$

or, in a thresholded implementation,

$$\max_{\pi \in \Pi^{\text{feas}}} \mathcal{W}_{\text{durable}}(\pi) \quad (6.2)$$

$$\text{subject to } \sup_{\theta \in \Theta_t} \mathbb{P}_{\theta}(\tau_{\chi_{\text{irr}}}^{(t)} \leq H \mid \pi) \leq \varepsilon, \quad (6.3)$$

$$R_t^J \geq J_{\text{min}}, \quad (6.4)$$

$$\mathcal{L}_{\text{nonfungible}}(\pi) = 0. \quad (6.5)$$

Species extinction, destruction of irreplaceable ecosystems, or intolerable distributional harm are not assumed tradable against aggregate income.

6.2 Robust barrier conditions

Let $h_{j,\theta}(\mathbf{x}) \geq 0$ encode safety condition j under model θ . A governance action is admissible only if

$$\sup_{\theta \in \Theta_t} \mathbb{P}_{\theta} \left[\min_j h_{j,\theta}(\mathbf{x}_{t+1}) < 0 \mid \mathbf{x}_t, \mathbf{u}_t \right] \leq \varepsilon_t. \quad (6.6)$$

A stronger margin-preservation condition is

$$\inf_{\theta \in \Theta_t} \min_j h_{j,\theta}(\mathbf{x}_{t+1}) \geq (1 - \kappa_t) \inf_{\theta \in \Theta_t} \min_j h_{j,\theta}(\mathbf{x}_t). \quad (6.7)$$

These equations are institutional design analogies to robust safety control, not claims that democratic governance admits exact barrier certificates.

6.3 Payoff inversion controller

The controller must change the utility differential, not merely issue a signal. The relevant chain is

$$\mathbf{k}_t \longrightarrow \mathbf{u}_{t+1} \longrightarrow \Delta_i(t+1) \longrightarrow \mathbf{a}_{i,t+1} \longrightarrow \mathbf{s}_{t+2}. \quad (6.8)$$

A sufficient local design target for the decisive coalition is

$$\begin{aligned} \Delta_i = & \underbrace{B_i^{\text{transition}} + B_i^{\text{market}} + B_i^{\text{legitimacy}}}_{\text{cooperation gains}} + \underbrace{C_i^{\text{avoided penalty}}}_{\text{credible enforcement}} \\ & - \underbrace{C_i^{\text{transition}}}_{\text{conversion burden}} - \underbrace{B_i^{\text{fossil}} + B_i^{\text{free}} + B_i^{\text{greenwash}} + B_i^{\text{delay}}}_{\text{defection gains}} > 0. \end{aligned} \quad (6.9)$$

The design must both raise the expected cost of defection and lower the feasible cost of cooperation. Penalty without transition support creates backlash; subsidy without enforcement preserves free-riding.

6.4 Trigger modes

Table 2: Reserve-based control modes. Thresholds must be domain-specific and pre-legislated.

Mode	Entry condition	Illustrative response
0: Adaptive normal	$\underline{CR}_t > m_1$	Monitoring, low-regret investment, adaptive pathway maintenance, routine redistribution.
1: Precaution	$0 < \underline{CR}_t \leq m_1$ or boundary uncertainty above tolerance	Stronger carbon and methane controls, tighter supply caps, higher audit frequency, stricter proof burden for one-way actions.
2: Emergency braking	$\underline{CR}_t \leq 0$ or high probability of near-term negative reserve	Temporary moratoria on new irreversible commitments, accelerated demand reduction, finance reallocation, rapid transition and justice buffers.
3: Loss minimization	$\rho_t(H) < \rho_{\min}$	Protect life-support systems, relocation, compensation, biodiversity refugia, cultural preservation, and safeguarding of still-viable subsystems.

Every trigger requires a release rule. Otherwise emergency governance can itself be captured and made permanent.

6.5 One-way doors and burden of proof

For normalized factors $I, S, D, U \in [0, 1]$ representing irreversibility, scale, duration, and uncertainty, define a One-Way Door index

$$\text{OWD}(a) = I(a) S(a) D(a) U(a), \quad (6.10)$$

As OWD rises, the acceptable risk $\varepsilon(a)$ falls. High-OWD actions are permitted only when the proponent demonstrates robust safety relative to the status quo:

$$\sup_{\theta \in \Theta_t} \mathbb{P}_\theta(\tau_{\chi_{\text{irr}}}^{(t)} \leq H \mid a) \leq \varepsilon(a), \quad \varepsilon'(\text{OWD}) < 0. \quad (6.11)$$

The status quo is not a zero-risk baseline. A new action may be precautionary when inaction has higher irreversible risk.

6.6 Accountable automation and institutional separation

Automation here means rule-bound transmission, not removal of politics. The policy rule is democratically authorized in advance; implementation becomes less discretionary after the trigger fires. An override is admissible only with

$$\text{Override} = \text{PublicReason} + \text{Expiry} + \text{IndependentReview} + \text{Compensation} + \text{AuditTrail}. \quad (6.12)$$

Appeal must not create free delay; disputed liabilities can be placed in escrow.

The architecture separates six functions:

$$\text{Sensor} \neq \text{Modeler} \neq \text{Trigger} \neq \text{JusticeAllocator} \neq \text{Auditor} \neq \text{Appeal}. \quad (6.13)$$

No single actor should control measurement, model selection, KPI computation, enforcement, audit, and exception review.




Figure 1: Closed-loop architecture. The red dashed paths represent strategic attacks; the lower control path represents anti-capture defenses.

7 Observation Architecture and Ecological Sensing

7.1 A dashboard of non-fungible signals

A practical KPI vector should contain physically and institutionally distinct signals rather than one composite score:

$$\mathbf{k}_t = \left(K_t^E, K_t^{\text{tail}}, K_t^{\text{infra}}, K_t^{\text{prod/finance}}, K_t^{\text{bio}}, K_t^J, K_t^{\text{capture}}, K_t^\tau \right). \quad (7.1)$$

Representative components are real emissions and budget consumption; threshold exceedance frequency; infrastructure failure duration; fossil-production and financial-backflow gaps; ecosystem integrity; justice burden; detected capture; and response delay. The dashboard is not aggregated when aggregation would allow one improvement to compensate for an irreplaceable loss.

7.2 Soundscape as one ensemble sensor

Soundscape ecology distinguishes biophony, geophony, and anthrophony and provides a scalable method for observing ecological change (Pijanowski et al., 2011). The Acoustic Complexity Index (ACI) measures

temporal changes in intensity within frequency bins (Pieretti et al., 2011). A schematic form is

$$\text{ACI} = \sum_f \frac{\sum_t |I_{f,t+1} - I_{f,t}|}{\sum_t I_{f,t} + \epsilon}. \quad (7.2)$$

The Normalized Difference Soundscape Index (NDSI) contrasts selected biophony and anthrophony bands:

$$\text{NDSI} = \frac{B - A}{B + A + \epsilon}. \quad (7.3)$$

Automated soundscape analysis has been used to track faunal recovery across tropical forest restoration gradients (Müller et al., 2023). However, meta-analysis shows that acoustic indices are context-dependent proxies rather than universal measures of biodiversity (Alcocer et al., 2022). Therefore

$$\widehat{B}_t = \mathcal{F}_{\text{eco}}(\text{soundscape, eDNA, remote sensing, field survey, camera traps, local knowledge}), \quad (7.4)$$

with uncertainty $\sigma_{B,t}$ and site-specific validation. The trigger should use posterior degradation probability, not raw loudness:

$$\mathbb{P}(B_t < B_{\text{baseline}} - \delta \mid \mathbf{y}_{0:t}) > p_{\text{crit}} \implies \text{review or response}. \quad (7.5)$$

Randomized sensor placement, hidden audit sensors, continuous sampling, playback anomaly detection, and protected access for sensitive species locations are part of the measurement design.

8 Red-Team: How the New OS Will Be Gamed

A governance model that assumes compliance after adoption is incomplete. Regulated actors adapt. Let the attack policy be

$$\pi_i^{\text{hack}} \in \arg \max_{m_i} \left[B_i^{\text{rent}} + B_i^{\text{delay}} + B_i^{\text{legitimacy}} + B_i^{\text{access}} - C_i^{\text{penalty}} - C_i^{\text{audit}} - C_i^{\text{legal}} \right]. \quad (8.1)$$

The target is often not physical improvement but maximal legal score at minimal transformation cost.

Table 3: Representative attack–defense pairs.

Attack surface	Realistic adaptation	Defense
Boundary laundering	Outsource or sell high-emission assets; reduce territorial or Scope 1 emissions while value-chain or consumption emissions rise.	Combine organizational, product, facility, consumption, and ownership accounting; retain transition responsibility after divestment for a defined period.
Offset substitution	Purchase low-cost credits and present finance as direct abatement.	Separate reduction, removal, and contribution claims; restrict credits to verified residual emissions; impose additionality, durability, and double-counting liability.
KPI cherry-picking	Improve carbon intensity while absolute emissions rise; emphasize renewable investment while fossil production expands.	Use mandatory KPI bundles, hard non-fungible floors, fixed core metrics plus rotating audit metrics, and absolute as well as intensity measures.

Attack surface	Realistic adaptation	Defense
Financial instrument gaming	Select easy sustainability targets, weak coupon penalties, or redemption before assessment.	Require material penalties, no early escape before target dates, independent baselines, absolute outcome metrics, and consequences beyond coupon adjustment.
Trade circumvention	Re-route goods, alter product classification, split shipments, or supply low-emission lots only to regulated markets.	Facility-level product passports, beneficial-ownership tracking, anti-circumvention rules, and conservative default factors when traceability is absent.
Soundscape theatre	Place sensors in favorable patches, shift disturbance outside sampling windows, or play recorded biophony.	Stratified randomized placement, continuous and surprise sampling, cross-sensor fusion, playback anomaly detection, and ecological composition metrics.
Adaptation substitution	Claim adaptation expenditure while mitigation declines or carbon lock-in grows.	Score adaptation net of induced emissions, lock-in, burden transfer, and maladaptation; do not allow adaptation to offset mitigation backsliding.
Justice backlash engineering	Concentrate transition costs on low-income households or dependent regions, then mobilize resulting opposition.	Couple every price or supply trigger to dividends, affordability floors, worker protection, regional investment, participation, and loss-and-damage finance.
Legal delay	Use review, litigation, transition periods, and exceptions to postpone enforcement while rents continue.	Automatic default rules, expiry, escrowed liabilities, time-limited appeals, and explicit delay costs.
Narrative capture	Reframe “transition” as indefinite preparation; use future technology promises to justify present expansion.	Evidence-backed claim standards, boundary disclosure, delivered-outcome accounting, legal liability, and zero credit for undelivered removals.
Data flooding	Publish voluminous, incomparable reports that increase audit cost and hide material information.	Machine-readable minimum schemas, open metadata, facility-level data, independent assurance, auditor rotation, and public replication.
Emergency-power capture	Use safety triggers to centralize authority indefinitely or suppress contestation.	Predefined release rules, judicial review, sunset clauses, rights floors, transparent models, and compensation for erroneous triggers.

The attack catalogue implies six general defenses:

$$\text{AntiCapture} = \text{Triangulation} + \text{Automation} + \text{Liability} + \text{AntiDelay} + \text{Justice} + \text{AdversarialAudit}. \quad (8.2)$$

The effective capture term may be represented as

$$c_t^{\text{eff}} = \frac{c_t^{\text{attempted}}}{1 + \text{AntiCapture}_t + R_t^J + \text{LegalExposure}_t}. \quad (8.3)$$

This expression is schematic; its empirical content would require domain-specific estimation.

9 Empirical Anchors and Claim Boundaries

Table 4: Selected empirical anchors. They motivate variables and failure modes but do not validate the full model.

Observation	Evidence	Model relevance
Non-punitive Paris compliance architecture	Article 15 and the PAICC specify a facilitative, non-adversarial, non-punitive mechanism (United Nations Framework Convention on Climate Change, 2015; United Nations Framework Convention on Climate Change, 2026).	Illustrates weak direct penalty coupling; does not imply that transparency and facilitation have zero value.
Persistent policy-temperature gap	UNEP 2025 estimated 2.3–2.5 °C under full NDC implementation and 2.8 °C under current policy (United Nations Environment Programme, 2025).	Motivates separating commitments from physical trajectory.
Fossil-production inconsistency	Governments collectively planned more than twice the 2030 production consistent with a 1.5 °C pathway (Stockholm Environment Institute et al., 2025).	Motivates a supply-side state and production-gap trigger.
Record atmospheric concentrations	WMO reported 423.9 ppm globally averaged CO ₂ in 2024 (World Meteorological Organization, 2025).	Anchors the physical state variable; concentration is not a complete welfare metric.
Multiple Earth-system boundaries	Recent boundary assessments integrate climate, biosphere, water, nutrients, and justice (Rockström et al., 2023; Richardson et al., 2023).	Supports coupled constraints rather than a single temperature KPI.
Context-sensitive acoustic indicators	Soundscape monitoring can track ecological recovery, while acoustic indices have limited universal proxy validity (Müller et al., 2023; Alcocer et al., 2022).	Supports ensemble sensing and explicit uncertainty rather than a single acoustic score.

The paper advances an architecture, not the following stronger claims:

- It does not prove that the Paris Agreement was intentionally designed to fail.
- It does not identify one exact global recovery boundary.
- It does not assume that every actor is selfish, identical, or perfectly rational.
- It does not imply that punishment alone is optimal, or that all cooperation should be centralized.

- It does not treat soundscape indices as a universal biodiversity truth machine.
- It does not show that an automatic controller can resolve distributive conflict without politics.

10 A Falsifiable Research Program

10.1 Minimal simulation environment

A first implementation should be an agent-based partially observable stochastic game, not a full integrated assessment model. Let actors vary in economic dependence, power, exposure, externalization capacity, horizon, trust, and capture ability. Compare at least four governance arms:

1. **Weak-coupling baseline:** observation and voluntary reporting with low $\partial \mathbf{t} / \partial \mathbf{k}$.
2. **Nominal trigger:** deterministic institutional response without capture defense.
3. **Robust reserve controller:** ambiguity-aware triggers using CR_t and viability constraints.
4. **Full RVCIM:** robust triggers plus justice buffers, model plurality, adversarial audit, and accountable override.

Key experimental sweeps include physical uncertainty, observation noise, common-mode model error, capture budget, policy delay, capital turnover, trust decay, and unequal transition burden. Outcomes include

$$R_{\text{irr}}(H) = \mathbb{P}(\tau_{X_{\text{irr}}}^{(t)} \leq H), \quad (10.1)$$

$$CR_{\text{min}} = \min_{t \leq H} CR_t^{\text{sys}}, \quad (10.2)$$

$$PG_{\text{cum}} = \sum_{t \leq H} PG_t, \quad (10.3)$$

$$CA = 1 - \frac{\|\mathbf{c}^{\text{eff}}\|}{\|\mathbf{c}^{\text{attempted}}\| + \epsilon}, \quad (10.4)$$

$$JS = R_t^J - \text{BacklashRisk}_t. \quad (10.5)$$

False-positive and false-negative trigger rates must both be measured. For irreversible risks, a false negative may have much greater cost, but excessive false positives can destroy legitimacy and thus future control capacity.

10.2 Hypotheses

Hypothesis 10.1 (Coupling). Holding physical dynamics, actor preferences, and capture fixed, increasing the credible KPI-to-payoff coupling Λ_i for decisive actors reduces defective action relative to a monitoring-only baseline.

Hypothesis 10.2 (Capture decay). A nominal trigger initially improves outcomes but its effect decays under repeated strategic adaptation when capture is unmodeled; adversarial audit and function separation reduce that decay.

Hypothesis 10.3 (Reserve advantage). Under delayed and uncertain tipping dynamics, reserve-based triggers reduce entry into X_{irr} relative to damage-threshold triggers, at the cost of a higher false-positive rate.

Hypothesis 10.4 (Justice durability). Transition policies with automatic justice buffers produce lower backlash, higher coordination reserve, and greater long-run policy persistence than otherwise identical

punitive policies.

Hypothesis 10.5 (Sensor ensemble). A multimodal ecological observation ensemble is more robust to manipulation and domain shift than a single KPI, but remains vulnerable to common-mode model error.

Hypothesis 10.6 (Adaptive path value). Policies with staged commitments, signposts, and reversible early actions retain more option value under deep uncertainty than fixed long-horizon plans with equivalent expected benefits.

10.3 Stress tests

At minimum, evaluate scenarios in which: the boundary is 30% closer than the median model; no reliable early warning appears; all models share a common bias; KPI data are accurate but institutions fail to act; institutions act but backlash forces repeal; an industry coordinates manipulation; two physical thresholds cascade; and emergency powers are captured. A design that succeeds only when all actors report honestly and all models are well calibrated is not an anti-capture architecture.

11 Executable Reference and Feasibility Contract

The repository companion turns the minimal research program into a small, inspectable experiment. Its purpose is to establish that the paper's formal pieces can be composed into one deterministic software loop and subjected to failure tests. Execution is an audit surface, not an evidentiary shortcut.

11.1 Feasibility ladder and admissible claims

The distinction between implementation and validation is organized as a four-level ladder. Progress at one level does not silently confer the claims of the next.

Table 5: Feasibility ladder for the manuscript and companion implementation.

Level	Object	Admissible claim	Additional evidence required to advance
F0	Executable structural toy	The declared mechanisms compose into a coherent, inspectable, replayable loop under recorded conditions.	Deterministic tests, explicit assumptions, traceable outputs, and hash-addressed receipts anchored by a trusted manifest.
F1	Stylized calibrated model	A narrow mechanism reproduces selected regularities under documented assumptions.	Parameter provenance, identification argument, alternative mechanisms, sensitivity analysis, and held-out tests.
F2	Sectoral or jurisdictional shadow model	The system may support bounded, non-actuating analysis in a declared domain.	Scoped legal authority, live-data validation, distribution-shift monitoring, independent review, and prospective evaluation.
F3	Reversible institutional pilot	A specific mechanism is operationally feasible under declared safeguards.	Rights floor, appeal, compensation, rollback and release rules, accountable ownership, and measured real-world outcomes.

This release is F0. A successful run, an ordered comparison table, or a lower toy irreversible-entry rate cannot be promoted into a climate forecast, a causal policy estimate, or evidence that an institutional

design is legitimate.

11.2 Paper-to-code correspondence

Table 6 makes the correspondence explicit enough to audit, while also stating where the implementation compresses the formal object.

Table 6: Correspondence between formal objects and the reference implementation.

Paper object	Implementation surface	F0 interpretation
x_t, s_t	WorldState	Normalized physical pressure, biosphere, institution, trust, justice, policy, support, and audit state.
$a_{i,t}, \Delta_i$	actor_actions()	Heterogeneous payoff-sensitive cooperation, defection, capture, restoration, and symbolic action.
y_t, k_t	observe_pressure()	Noisy observation with manipulation and audit correction.
Θ_t	Environment.model_boundaries	A finite model ensemble with common-mode bias and structural allowance.
CR_t, ρ_t	estimate_reserve()	Conservative estimated exit time minus endogenous response time.
Γ	select_mode(), schedule_policy(), advance_policy()	Trigger choice, delayed activation, and release logic.
c_t^{eff}	effective_capture_value()	Attempted capture after audit, function separation, justice, and policy exposure.
PG_t	Trace and episode accumulators	Symbolic action minus scaled verified physical improvement.

The code compares the four arms specified in the research program: weak coupling, nominal trigger, robust reserve, and full RVCIM. Within an episode, all arms receive the same actor profiles, realized boundary, model errors, disturbances, observation noise, social noise, and random choice draws. This paired construction reduces internal comparison noise; it does not establish external validity.

11.3 Reference implementation contract

The executable contract is intentionally narrow and fail-closed:

1. **Runtime.** The simulator uses Python 3.10 or newer and only the standard library. The checked-in source is the readable source of truth; execution does not depend on hidden payload assembly.
2. **Declared input.** A run resolves simulation/configs/minimal.json, selected arms, episode count, seed, and command-line overrides into resolved_config.json. Unknown configuration keys are rejected.
3. **Paired randomness.** The episode seed determines one common environment for all selected arms. Arm ordering cannot change the sampled exogenous episode.
4. **Observable outputs.** Every complete run writes summary.csv, episodes.csv, trace.csv, comparison.md, resolved_config.json, and receipt.json. The receipt records the claim boundary, command, environment, source identity, configuration identity, and output hashes.
5. **Verification.** Receipt verification resolves recorded paths relative to the receipt, recomputes SHA-256

identities, and returns non-zero for a missing or modified file. Repository verification also checks the release manifest and committed reference receipt, then replays the declared reference command and byte-compares its deterministic outputs.

6. **Failure visibility.** False-positive and false-negative trigger rates are reported together. Skipped executable tests, absent release artifacts, and hash mismatches are failures, not successful partial results. Manuscript buildability is qualified separately with an explicit TeX toolchain and must not be inferred from Python CI.

The reference command is:

```
python3 -m simulation run \
  --config simulation/configs/minimal.json \
  --episodes 64 --seed 7 \
  --out artifacts/reference_run --overwrite

python3 -m simulation verify \
  --receipt artifacts/reference_run/receipt.json
```

Claim boundary

What F0 does not establish. The normalized parameters are selected to expose an inspectable mechanism gradient. They do not estimate climate sensitivity, tipping distributions, institutional response, actor preferences, or real intervention effects. The executable establishes correspondence, replayability under the recorded source, configuration, seed, and compatible runtime, and falsifiable interfaces—not empirical truth, policy rank, safety, legality, or democratic legitimacy.

12 Limitations and Open Problems

First, the cooperative/defective partition is partly normative. An action can reduce emissions while violating justice or biodiversity constraints. The paper therefore requires plural hard constraints rather than a single climate score, but disagreements over those constraints remain political.

Second, robust worst-case methods can become excessively conservative. Ambiguity sets may contain implausible models, and powerful actors may weaponize uncertainty to block beneficial transformation. The response is not to ignore uncertainty but to combine hard safety constraints with safe learning, staged trials, explicit information value, and periodic model-set revision.

Third, institutional actuators do not behave like physical actuators. Enforcement depends on legitimacy, legal authority, implementation capacity, and interpretation. Equations such as Equations (3.12) and (6.6) are formal abstractions that must be instantiated separately in each jurisdiction.

Fourth, automatic triggers create constitutional risk. Rules justified by an unknown boundary can be used to normalize emergency powers. Release conditions, rights floors, transparent models, independent appeal, compensation, and legislative renewal are therefore part of the controller, not optional safeguards.

Fifth, the robust viability kernel may be computationally intractable in high-dimensional coupled systems. Approximations, surrogate models, scenario discovery, and decomposition will be necessary. Model plurality can reduce single-model dependence but does not eliminate correlated error.

Sixth, physical stabilization after payoff inversion may be slow. A governance reform can be successful in strategic terms while concentrations, ice loss, or ecological damage continue due to lag. Evaluation must distinguish policy coupling, action change, and delayed physical response.

Finally, this paper has not estimated its parameters or demonstrated empirical superiority. It defines the variables, conditional claims, attack surfaces, and tests needed to do so.

13 Conclusion

Nash's Cage names a system in which local rationality, physical externalization, short horizons, partial observation, and institutional capture jointly stabilize collective degradation. The architecture is dangerous not only because the physical state worsens, but because it can shrink the future set of feasible recoveries. The decisive threshold may therefore precede visible collapse: remaining safe time becomes shorter than institutional and biophysical braking time.

RVCIM responds by changing the order of operations. It does not maximize short-term welfare and then price the damage. It first constrains irreversible risk and injustice, then optimizes within the remaining viable region. It does not assume that a better dashboard is governance. It measures the full chain from observation to actuator, payoff, behavior, and physical response. It does not assume reform survives contact with regulated actors. Capture is a strategic state variable, and anti-capture design is an institutional immune system.

The manuscript's core proposition is deliberately conservative:

Core claim

When the irreversible boundary is unknown, governance should preserve the ability to brake rather than infer safety from the absence of visible collapse.

Observation without actuation is ritual. Actuation without justice is unstable. Optimization without preserved controllability is recklessness. The research question is therefore not simply whether the world is currently inside a nominal target. It is whether, across plausible models and real institutional delays, a feasible path back to a safe and just region still exists – and whether the rules of the game make taking that path rational before the reserve is gone.

A Notation Summary

Symbol	Meaning
x_t	Full physical, agent, institutional, and narrative state.
s_t	Physical and ecological state.
$a_{i,t}$	Actor i 's emissions, fossil, decarbonization, symbolic, capture, restoration, and adaptation actions.
ι_t	Institutional actuators and capacities.
n_t	Narrative, awareness, trust, and legitimacy state.
y_t	Raw multimodal observations.
k_t	Institutional KPI and reserve dashboard.
c_t	Capture vector attacking metrics, institutions, delay, narrative, justice, and finance.
Θ_t	Ambiguity set of plausible system models.
$\mathcal{X}_{\text{safe}}$	Present-time safe and just constraint set.
$\mathcal{V}_t^{\text{rob}}$	Robust viability set under feasible common policies.
$\mathcal{X}_{\text{irr}}$	Functionally irreversible state set.
$\tau_{\mathcal{A}}^{(t)}$	First-passage delay from decision time t into set $\mathcal{A}$.
$\rho_t(H)$	Robust recoverability over horizon H .

CR_t	Remaining lower-bound exit time minus upper-bound response time.
Δ_i	Cooperation payoff minus defection payoff.
Λ_i	KPI-to-payoff coupling strength.
PG_t	Symbolic action minus scaled verified physical improvement.
R_t^J	Justice and legitimacy reserve.
$OWD(a)$	One-Way Door index for action a .

B Minimal Control Loop

At each decision time t :

1. Observe:
 $y_t \leftarrow$ independent multi-sensor observations
2. Update uncertainty:
 $\Theta_t \leftarrow$ update plausible models + structural-error allowance
3. Estimate state and reserves:
 $k_t, V_{rob_t}, \rho_t, CR_t \leftarrow$ robust inference($y_{0:t}, \Theta_t$)
4. Red-team the observation:
 $k_{robust} \leftarrow$ sensor fusion - manipulation allowance + audit correction
5. Select mode:
 if $\rho_t < \rho_{min}$:
 mode $\leftarrow$ loss minimization
 else if $lower_bound(CR_t) \leq 0$:
 mode $\leftarrow$ emergency braking
 else if CR_t is low or boundary uncertainty is high:
 mode $\leftarrow$ precaution
 else:
 mode $\leftarrow$ adaptive normal
6. Select policy:
 maximize durable welfare + information value + option value
 subject to:
 irreversible-risk $\leq \epsilon$
 justice reserve $\geq J_{min}$
 prohibited non-fungible loss $== 0$
7. Trigger actuators:
 price, cap, permit, finance, restoration, dividend, compensation
8. Audit adaptation:
 search for metric, boundary, delay, narrative, finance, and justice hacks
9. Publish:
 observations, models, uncertainty, trigger, burdens, beneficiaries,
 override and release conditions
10. Require $review_interval < controllability_reserve$; otherwise increase review cadence.

C Cage-Inversion Constitutional Invariants

1. **Unknown is not zero.** Missing evidence is not a safety certificate.
2. **No silent loss of controllability.** Reserve deterioration must itself be a trigger.
3. **No compensation across irreplaceable boundaries.** Non-fungible losses cannot be converted into generic credits.
4. **One-way actions carry the proof burden.** Irreversibility raises the required confidence of safety.
5. **Delay has a price.** Appeal and review cannot be costless defection strategies.
6. **Justice is a floor.** Aggregate benefit cannot erase intolerable burden concentration.
7. **No single epistemic controller.** Measurement, modeling, triggering, auditing, and appeal are separated.
8. **Override is temporary and expensive.** Exceptions require public reasons, expiry, review, compensation, and traceability.
9. **Every emergency trigger has a release rule.** Safety governance cannot become an unbounded permanent exception.

References

- Alcocer, Irene, Herlander Lima, Larissa S. M. Sugai, and Diego Llusia (2022). “Acoustic Indices as Proxies for Biodiversity: A Meta-Analysis”. In: *Biological Reviews* 97.6, pp. 2209–2236. doi: 10.1111/brv.12890.
- Ames, Aaron D., Samuel Coogan, Magnus Egerstedt, Gennaro Notomista, Koushil Sreenath, and Paulo Tabuada (2019). “Control Barrier Functions: Theory and Applications”. In: *2019 18th European Control Conference*, pp. 3420–3431. doi: 10.23919/ECC.2019.8796030.
- Ames, Aaron D., Xiangru Xu, Jessy W. Grizzle, and Paulo Tabuada (2017). “Control Barrier Function Based Quadratic Programs for Safety Critical Systems”. In: *IEEE Transactions on Automatic Control* 62.8, pp. 3861–3876. doi: 10.1109/TAC.2016.2638961.
- Armstrong McKay, David I., Arie Staal, Jesse F. Abrams, Ricarda Winkelmann, Boris Sakschewski, Sina Loriani, Ingo Fetzer, Sarah E. Cornell, Johan Rockström, and Timothy M. Lenton (2022). “Exceeding 1.5 °C Global Warming Could Trigger Multiple Climate Tipping Points”. In: *Science* 377.6611, eabn7950. doi: 10.1126/science.abn7950.
- Aubin, Jean-Pierre, Alexandre M. Bayen, and Patrick Saint-Pierre (2011). *Viability Theory: New Directions*. Berlin: Springer. doi: 10.1007/978-3-642-16684-6.
- Axelrod, Robert (1984). *The Evolution of Cooperation*. New York: Basic Books.
- Başar, Tamer and Geert Jan Olsder (1999). *Dynamic Noncooperative Game Theory*. 2nd ed. Philadelphia: Society for Industrial and Applied Mathematics. doi: 10.1137/1.9781611971132.
- Carpenter, Daniel and David A. Moss, eds. (2014). *Preventing Regulatory Capture: Special Interest Influence and How to Limit It*. Cambridge: Cambridge University Press. doi: 10.1017/CBO9781139565875.
- Delmas, Magali A. and Vanessa Cuerel Burbano (2011). “The Drivers of Greenwashing”. In: *California Management Review* 54.1, pp. 64–87. doi: 10.1525/cmr.2011.54.1.64.
- Haasnoot, Marjolijn, Jan H. Kwakkel, Warren E. Walker, and Judith ter Maat (2013). “Dynamic Adaptive Policy Pathways: A Method for Crafting Robust Decisions for a Deeply Uncertain World”. In: *Global Environmental Change* 23.2, pp. 485–498. doi: 10.1016/j.gloenvcha.2012.12.006.
- Hardin, Garrett (1968). “The Tragedy of the Commons”. In: *Science* 162.3859, pp. 1243–1248. doi: 10.1126/science.162.3859.1243.
- Intergovernmental Panel on Climate Change (2022). *Climate Change 2022: Impacts, Adaptation and Vulnerability*. Cambridge University Press. doi: 10.1017/9781009325844.
- (2023). *Climate Change 2023: Synthesis Report*. Geneva: IPCC. doi: 10.59327/IPCC/AR6-9789291691647.

- Kaelbling, Leslie Pack, Michael L. Littman, and Anthony R. Cassandra (1998). "Planning and Acting in Partially Observable Stochastic Domains". In: *Artificial Intelligence* 101.1–2, pp. 99–134. doi: 10.1016/S0004-3702(98)00023-X.
- Kwakkel, Jan H., Marjolijn Haasnoot, and Warren E. Walker (2016). "Comparing Robust Decision-Making and Dynamic Adaptive Policy Pathways for Model-Based Decision Support under Deep Uncertainty". In: *Environmental Modelling & Software* 86, pp. 168–183. doi: 10.1016/j.envsoft.2016.09.017.
- Milinski, Manfred, Ralf D. Sommerfeld, Hans-Juergen Krambeck, Floyd A. Reed, and Jochem Marotzke (2008). "The Collective-Risk Social Dilemma and the Prevention of Simulated Dangerous Climate Change". In: *Proceedings of the National Academy of Sciences* 105.7, pp. 2291–2294. doi: 10.1073/pnas.0709546105.
- Müller, Jörg, Oliver Mitesser, H. Martin Schaefer, Sebastian Seibold, Adrian Busse, Peter Kriegel, Dominik Rabl, Roberto Gelis, Andrea Arteaga, Juan Freile, et al. (2023). "Soundscapes and Deep Learning Enable Tracking Biodiversity Recovery in Tropical Forests". In: *Nature Communications* 14, p. 6191. doi: 10.1038/s41467-023-41693-w.
- Nash, John F. Jr. (1950). "Equilibrium Points in n -Person Games". In: *Proceedings of the National Academy of Sciences* 36.1, pp. 48–49. doi: 10.1073/pnas.36.1.48.
- Nordhaus, William (2015). "Climate Clubs: Overcoming Free-Riding in International Climate Policy". In: *American Economic Review* 105.4, pp. 1339–1370. doi: 10.1257/aer.15000001.
- Ostrom, Elinor (1990). *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge: Cambridge University Press. doi: 10.1017/CBO9780511807763.
- Pieretti, Nadia, Almo Farina, and Davide Morri (2011). "A New Methodology to Infer the Singing Activity of an Avian Community: The Acoustic Complexity Index". In: *Ecological Indicators* 11.3, pp. 868–873. doi: 10.1016/j.ecolind.2010.11.005.
- Pijanowski, Bryan C., Almo Farina, Stuart H. Gage, Sarah L. Dumyahn, and Bernie L. Krause (2011). "Soundscape Ecology: The Science of Sound in the Landscape". In: *BioScience* 61.3, pp. 203–216. doi: 10.1525/bio.2011.61.3.6.
- Richardson, Katherine, Will Steffen, Wolfgang Lucht, Jørgen Bendtsen, Sarah E. Cornell, Jonathan F. Donges, Markus Drüke, Ingo Fetzer, Govindasamy Bala, Werner von Bloh, et al. (2023). "Earth beyond Six of Nine Planetary Boundaries". In: *Science Advances* 9.37, eadh2458. doi: 10.1126/sciadv.adh2458.
- Rockström, Johan, Joyeeta Gupta, Dahe Qin, Steven J. Lade, Jesse F. Abrams, Lauren S. Andersen, David I. Armstrong McKay, Xuemei Bai, Govindasamy Bala, Stuart E. Bunn, et al. (2023). "Safe and Just Earth System Boundaries". In: *Nature* 619, pp. 102–111. doi: 10.1038/s41586-023-06083-8.
- Scheffer, Marten, Jordi Bascompte, William A. Brock, Victor Brovkin, Stephen R. Carpenter, Vasilis Dakos, Hermann Held, Egbert H. van Nes, Max Rietkerk, and George Sugihara (2009). "Early-Warning Signals for Critical Transitions". In: *Nature* 461, pp. 53–59. doi: 10.1038/nature08227.
- Stockholm Environment Institute, Climate Analytics, E3G, International Institute for Sustainable Development, and United Nations Environment Programme (2025). *The Production Gap Report 2025*. Production Gap Report Initiative. URL: <https://productiongap.org/2025report/> (visited on 07/30/2026).
- United Nations Environment Programme (2025). *Emissions Gap Report 2025*. Nairobi: UNEP. URL: <https://www.unep.org/resources/emissions-gap-report-2025> (visited on 07/30/2026).
- United Nations Framework Convention on Climate Change (2015). *Paris Agreement*. United Nations Treaty Series. URL: https://unfccc.int/sites/default/files/english_paris_agreement.pdf (visited on 07/30/2026).
- (2026). *More about the Paris Agreement Implementation and Compliance Committee*. URL: <https://unfccc.int/process-and-meetings/bodies/constituted-bodies/paris-agreement-implementation-and-compliance-committee-paicc/more-about-paicc> (visited on 07/30/2026).
- World Meteorological Organization (2025). *WMO Greenhouse Gas Bulletin No. 21: The State of Greenhouse Gases in the Atmosphere Based on Global Observations through 2024*. Geneva: WMO. URL: <https://wmo.int/resources/publication-series/greenhouse-gas-bulletin/wmo-greenhouse-gas-bulletin-no-21> (visited on 07/30/2026).